

Fast Option Pricing using Fourier Transform

When Monte-Carlo is too slow



VERTOX

JUN 09, 2026



19



2



5

Share

Monte-Carlo Simulation is the most straightforward way to price an option, and if you don't care about speed, it's a solid choice.

The moment you care about speed, like when quoting live, or when calibrating a pricing model where you need to reprice options thousands of times, Monte Carlo quickly becomes unusable.

This article presents three Fourier-based methods that solve this problem, reducing pricing from seconds to microseconds. We'll use the Heston model as our running example throughout, and achieve a speedup of 30,000x against Monte Carlo!

I write about quantitative trading the way it's actually practised:

Robust models and portfolios, combining signals and strategies, understanding the assumptions behind your models.

Topics I write about include portfolio construction, market making, risk management, research methodology, and more.

If this way of thinking resonates, you'll probably like what I publish.

VertoxQuant is a reader-supported publication.

To receive new posts and support my work,

consider becoming a free or paid subscriber.

What you'll learn

- Why Monte Carlo breaks down for live pricing and model calibration, and what makes it fundamentally unsuitable for production options trading.
- How the Fourier Transform and FFT work and why they matter for options pricing.
- What the characteristic function is and how it enables us to price options even when the risk-neutral density of the model has no known closed-form analytical solution.
- How Carr-Madan transforms option pricing into an FFT problem.
- How Lewis reformulates the same problem via contour integration, eliminating the dampening parameter that makes Carr-Madan brittle.
- How the COS method takes a completely different approach, approximating the density via a cosine series, and why it converges exponentially fast.
- Python implementations for all methods, together with speed comparisons.

This article is free to read, so I would appreciate it if you shared it with someone who would also like it!

The Characteristic Function

To price options using Fourier methods, we need one thing from our model: the characteristic function of $\log S_T$ where S_T is the price at expiry.

For a random variable X , the characteristic function is defined as:

$$\phi(u) = \mathbb{E}[e^{iuX}]$$

For $\log S_T$ we write:

$$\phi_T(u) = \mathbb{E}^Q[e^{iu \log S_T}],$$

where Q is the risk-neutral measure.

Just like the probability density function, it completely characterises the distribution of X . But while the density $q(x)$ often has no closed-form expression, the characteristic function often does and is much easier to find.

In fact, the characteristic function and the density are a Fourier transform pair. Given the characteristic function, you can recover the density via the inverse Fourier transform:

$$q(x) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-iux} \phi_T(u) du$$

This is the core idea behind all Fourier pricing methods. You don't need to know the distribution of $\log S_T$ explicitly; you just need the characteristic function.

Why does the characteristic function always exist?

For any model where $\log S_T$ is a well-defined random variable, the characteristic function is guaranteed to exist for all real u , since

$$|e^{iu \log S_T}| = 1,$$

and thus the expectation is always finite.

Pricing in Fourier Space

In theory, you could now go ahead and compute the distribution from the characteristic function numerically, and then integrate against the payoff to get the

fair value of the option. In practice, Fourier pricing methods skip this step entirely price directly in Fourier space, which is both faster and more numerically stable.

Two Examples

For Black-Scholes, $\log S_T$ is Gaussian, so the characteristic function is that of a Gaussian random variable, with appropriate mean and standard deviation plugged in

$$\phi_T^{\text{BS}}(u) = \exp(iu(\log S_0 + (r - \frac{\sigma^2}{2})T) - \frac{u^2 \sigma^2 T}{2})$$

For Heston, $\log S_T$ has no nice closed-form density expression. But the characteristic function is still available in closed form:

$$\phi_T^{\text{Heston}}(u) = \exp(iu \log S_0 + A(T, u) + B(T, u)v_0)$$

where:

$$A(T, u) = iurT + \frac{\kappa\theta}{\xi^2} [(\rho\xi iu - \kappa - d)T - 2 \log \frac{1 - ge^{-dT}}{1 - g}]$$

$$B(T, u) = \frac{(\rho\xi iu - \kappa - d)(1 - e^{-dT})}{\xi^2(1 - ge^{-dT})}$$

with

$$d = \sqrt{(\rho\xi iu - \kappa)^2 + \xi^2(iu + u^2)}$$

$$g = \frac{\rho\xi iu - \kappa - d}{\rho\xi iu - \kappa + d}$$

These are derived by applying Itô's lemma to $\log S_t$, exploiting the affine structure of the Heston dynamics, and solving the resulting Riccati ODEs in closed form.

How to properly derive this will be the subject of future articles.

Fourier Transform and FFT

Before deriving any pricing formulas, we need to understand the Fourier Transform and the Fast Fourier Transform algorithm.

The Fourier Transform

For a real square integrable function $f(x)$, the Fourier transform is:

$$\hat{f}(\xi) = \int_{-\infty}^{\infty} e^{i\xi x} f(x) dx$$

This function lives in frequency space; It tells you how much of each frequency ξ is present in f . You recover f from $\hat{f}$ via the inverse Fourier transform:

$$f(x) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-i\xi x} \hat{f}(\xi) d\xi.$$

Two real functions f and g that are both square integrable have a well-defined inner product. Parseval's theorem says this inner product can also be computed in Fourier space:

$$\int_{-\infty}^{\infty} f(x) \overline{g(x)} dx = \frac{1}{2\pi} \int_{-\infty}^{\infty} \hat{f}(\xi) \overline{\hat{g}(\xi)} d\xi$$

This will be used later on in the Lewis method.

The Discrete Fourier Transform

In practice, you work with a finite grid of N points rather than a continuous function. The Discrete Fourier Transform (DFT) of a sequence $x_0, x_1, \dots, x_{N-1}$ is:

$$X_k = \sum_{j=0}^{N-1} x_j e^{-i2\pi \frac{k}{N} j}.$$

Each output X_k is the inner product of the input sequence against a complex exponential at frequency k . Computing all N outputs naively requires evaluating N terms for each of the N outputs, a total of N^2 multiplications.

The Fast Fourier Transform

For N a power of 2, the Fast Fourier Transform, introduced by Cooley and Tukey, reduces the cost from $O(N^2)$ to $O(N \log N)$ by exploiting the fact that the DFT of a sequence of length N can be written as the sum of two DFTs of length $N/2$.

Applying this decomposition recursively, splitting each half into its own two halves reduces the total number of multiplications to $O(N \log N)$. For $N = 4096$, this is roughly a 340x speedup over the naive approach!

More formally, the DFT of a sequence of length N can be written as the sum of the DFT over the even-indexed elements and the DFT over the odd-indexed elements:

$$\begin{aligned} X_k &= \sum_{j=0}^{N/2-1} x_{2j} e^{-\frac{2\pi i}{N} (2j)k} + \sum_{j=0}^{N/2-1} x_{2j+1} e^{-\frac{2\pi i}{N} (2j+1)k} \\ &= E_k + e^{-\frac{2\pi i}{N} k} O_k \end{aligned}$$

where E_k and O_k are the DFTs of the even and odd subsequences.

Carr-Madan

The Carr-Madan method (1999) is one of the most popular option pricing methods using the Fast Fourier Transform.

Let $s = \log S_T$ and $k = \log K$. The risk-neutral density of s is $q_T(s)$, with characteristic function

$$\phi_T(u) = \int_{-\infty}^{\infty} e^{ius} q_T(s) ds.$$

The call price as a function of log strike is:

$$C_T(k) = e^{-rT} \int_k^{\infty} (e^s - e^k) q_T(s) ds$$

We want to Fourier transform $C_T(k)$. The problem is that $C_T(k)$ is not square-integrable, since as $k \rightarrow -\infty$, $C_T(k) \rightarrow S_0$ rather than decaying to zero. So its Fourier Transform doesn't exist in the classical sense.

The fix is to multiply by an exponential dampening factor. Define the modified call price:

$$c_T(k) = e^{\alpha k} C_T(k), \quad \alpha > 0.$$

For suitable α , $c_T(k)$ is square integrable, since the $e^{\alpha k}$ term decays fast enough as $k \rightarrow -\infty$ to kill the blow-up. Therefore, this has a well-defined Fourier transform:

$$\Psi_T(v) = \int_{-\infty}^{\infty} e^{ivk} c_T(k) dk$$

Our goal now is to write an analytical expression of Ψ_T in terms of ϕ_T , and then obtain call prices numerically using the inverse transform.

Deriving Ψ_T

Substitute the definitions of $c_T(k)$ and $C_T(k)$:

$$\Psi_T(v) = \int_{-\infty}^{\infty} e^{ivk} e^{\alpha k} \int_k^{\infty} e^{-rT} (e^s - e^k) q_T(s) ds dk$$

The region of integration is

$$\{(k, s) : k \in (-\infty, \infty), s \in (k, \infty)\},$$

which is equivalent to:

$$\{(k, s) : s \in (-\infty, \infty), k \in (-\infty, s)\}.$$

The integrand is non-negative, so we can apply Tonelli's theorem and swap the order of integration:

$$\Psi_T(v) = \int_{-\infty}^{\infty} e^{-rT} q_T(s) \int_{-\infty}^s e^{\alpha k} (e^s - e^k) e^{ivk} dk ds$$

Expand the inner integrand:

$$e^{\alpha k} (e^s - e^k) e^{ivk} = e^s e^{(\alpha+iv)k} - e^{(1+\alpha+iv)k}$$

For the first term:

$$e^s \int_{-\infty}^s e^{(\alpha+iv)k} dk = e^s \left[\frac{e^{(\alpha+iv)k}}{\alpha+iv} \right]_{-\infty}^s = \frac{e^{(\alpha+1+iv)s}}{\alpha+iv}$$

For the second term:

$$\int_{-\infty}^s e^{(1+\alpha+iv)k} dk = \left[\frac{e^{(1+\alpha+iv)k}}{1+\alpha+iv} \right]_{-\infty}^s = \frac{e^{(1+\alpha+iv)s}}{1+\alpha+iv}$$

Substituting back:

$$\Psi_T(v) = e^{-rT} \int_{-\infty}^{\infty} q_T(s) \left(\frac{e^{(\alpha+1+iv)s}}{\alpha+iv} - \frac{e^{(\alpha+1+iv)s}}{\alpha+1+iv} \right) ds$$

Factor out

$$e^{(\alpha+1+iv)s}$$

and recognise that since

$$e^{(\alpha+1+iv)s} = e^{i(v-(\alpha+1))s},$$

we have:

$$\int_{-\infty}^{\infty} q_T(s) e^{(\alpha+1+iv)s} ds = \phi_T(v - (\alpha + 1)i)$$

So:

$$\Psi_T(v) = e^{-rT} \phi_T(v - (\alpha + 1)i) \left(\frac{1}{\alpha+iv} - \frac{1}{\alpha+1+iv} \right)$$

After simplifying, we finally get:

$$\Psi_T(v) = \frac{e^{-rT} \phi_T(v - (\alpha+1)i)}{\alpha^2 + \alpha - v^2 + i(2\alpha+1)v}$$

This is expressed entirely in terms of ϕ_T , the characteristic function of $\log S_T$ under our model of choice, which we said earlier we often have a closed-form solution for, so this whole thing is closed-form!

Recovery Formula

Since Ψ_T is the Fourier transform of the dampened call price, standard Fourier inversion gives:

$$c_T(k) = \frac{1}{2\pi} \int_{-\infty}^{\infty} e^{-ivk} \Psi_T(v) dv$$

Dividing both sides by $e^{\alpha k}$:

$$C_T(k) = \frac{e^{-\alpha k}}{2\pi} \int_{-\infty}^{\infty} e^{-ivk} \Psi_T(v) dv$$

Since $C_T(k)$ is real, we have

$$\Psi_T(-v) = \overline{\Psi_T(v)},$$

so the integrand over $(-\infty, 0)$ is the complex conjugate of the integrand over $(0, \infty)$ and their sum is twice the real part:

$$C_T(k) = \frac{e^{-\alpha k}}{\pi} \int_0^{\infty} \Re[e^{-ivk} \Psi_T(v)] dv$$

Choice of alpha

Two conditions constrain α . First, for $\Psi_T(0)$ to be finite, we need

$$\phi_T(-(\alpha+1)i) = \mathbb{E}^{\mathbb{Q}}[S_T^{\alpha+1}] < \infty.$$

Another condition is α not being equal to 0, to avoid a singularity in the denominator at $v = 0$. For Black-Scholes and Heston, all positive moments exist, so

alpha is unconstrained. Carr and Madan suggest using one quarter of the upper bound implied by the moment condition, and find alpha = 1.5 works well in practice.

Truncation bound

The recovery integral runs from 0 to infinity. We need to verify it converges and can safely truncated at some finite value a. For that, let's first look at the numerator of $|\Psi_T(v)|$.

It's easy to see that

$$|\phi_T(v - (\alpha + 1)i)| \leq \mathbb{E}^{\mathbb{Q}}[S_T^{\alpha+1}],$$

which is independent of v.

The denominator of $\Psi_T(v)$ grows like v^4 for large v, so:

$$|\Psi_T(v)|^2 \leq \frac{\mathbb{E}[S_T^{\alpha+1}]}{(\alpha^2 + \alpha - v^2)^2 + (2\alpha + 1)^2 v^2} \leq \frac{A}{v^4}$$

for some constant A. Therefore

$$|\Psi_T(v)| \leq \frac{\sqrt{A}}{v^2}$$

and the tail integral from any truncation point a satisfies:

$$\int_a^\infty |\Psi_T(v)| dv \leq \frac{\sqrt{A}}{a}$$

Thus, the truncation error is bounded by:

$$|C_T(k) - \tilde{C}_T(k)| \leq \frac{e^{-\alpha k}}{\pi} \frac{\sqrt{A}}{a}$$

To achieve a target error epsilon, choose:

$$a > \frac{e^{-\alpha k} \sqrt{A}}{\pi \epsilon}$$

FFT discretization

Discretize the recovery integral with N points at spacing η , setting

$$v_j = \eta(j-1) \quad \text{for } j = 1, \dots, N$$

Applying the trapezoid rule:

$$C_T(k) \approx \frac{e^{-\alpha k}}{\pi} \sum_{j=1}^N e^{-iv_j k} \Psi_T(v_j) \eta$$

The effective upper limit of integration is $a = N^* \eta$. We want to evaluate this at a grid of N log-strikes

$$k_u = -b + \lambda(u-1) \quad \text{for } u = 1, \dots, N,$$

which gives us log strike levels ranging from $-b$ to b , centred around ATM by setting

$$b = \frac{1}{2} N \lambda.$$

Substituting v_j and k_u :

$$e^{-iv_j k_u} = e^{ibv_j} e^{-i\eta\lambda(j-1)(u-1)}$$

For the sum to match the DFT structure

$$e^{-i \frac{2\pi}{N} (j-1)(u-1)}$$

we need

$$\lambda\eta = \frac{2\pi}{N}.$$

This is the grid coupling constraint; Fine integration grid (η small) forces coarse strike spacing (λ large) and vice versa. The sum becomes:

$$C_T(k_u) \approx \frac{e^{-\alpha k_u}}{\pi} \sum_{j=1}^N e^{-i \frac{2\pi}{N} (j-1)(u-1)} e^{ibv_j} \Psi_T(v_j) \eta$$

which is exactly the DFT of the sequence

$$e^{ibv_j} \Psi_t(v_j) \eta.$$

A single FFT call evaluates this for all N strikes simultaneously.

So you build the input vector

$$x_j = e^{ibv_j} \Psi_T(v_j) \eta$$

once and then call the FFT algorithm. Then multiply the u-th output by

$$\frac{e^{-\alpha k_u}}{\pi}$$

to get the final call prices for each strike.

Simpson's rule

The trapezoid rule has errors $O(\eta^2)$. To get accurate integration with larger η , which, via the grid coupling constraint, means finer strike spacing, we replace the uniform weight η with Simpson's rule weights:

$$w_j = \frac{\eta}{3} [3 + (-1)^j + \delta_{j-1}]$$

where δ_{j-1} is the Kronecker delta, equal to 1 only at $j=1$. The error improves from $O(\eta^2)$ to $O(\eta^4)$, and the FFT structure is preserved since replacing η with w_j just multiplies each term by a scalar before applying the FFT.

The final formula is:

$$C_T(k_u) \approx \frac{e^{-\alpha k_u}}{\pi} \sum_{j=1}^N e^{-i \frac{2\pi}{N} (j-1)(u-1)} e^{ibv_j} \Psi_T(v_j) w_j$$

The weakness of Carr-Madan

The grid coupling constraint is the fundamental limitation; you cannot independently choose integration accuracy and strike grid density. Additionally, alpha requires model-specific tuning. Both issues are addressed by the methods that follow.

Python Implementation

We will be working with Heston in this article. Let's first implement Monte Carlo pricing using Euler-Maruyama as a ground truth and for speed comparison:

```
1 import numpy as np
2 from scipy.stats import norm
3 import time
4 import matplotlib.pyplot as plt
5 from scipy.integrate import quad
```

```
1 def heston_mc(S0, K, r, q, T, kappa, theta, eta, rho, v0,
2               n_steps=200, n_paths=1_000_000, seed=42):
3     rng = np.random.default_rng(seed)
4     dt = T / n_steps
5     S = np.full(n_paths, float(S0))
6     v = np.full(n_paths, float(v0))
7
8     for _ in range(n_steps):
9         Z1 = rng.standard_normal(n_paths)
10        Z2 = rho * Z1 + np.sqrt(1 - rho**2) * rng.standard_normal(n_paths)
11        v_pos = np.maximum(v, 0)
12        S *= np.exp((r - q - 0.5 * v_pos) * dt + np.sqrt(v_pos * dt) * Z1)
13        v += kappa * (theta - v_pos) * dt + eta * np.sqrt(v_pos * dt) * Z2
14
15    return np.exp(-r * T) * np.mean(np.maximum(S - K, 0))
```

Next, we implement the Heston Characteristic Function:

```
1 def heston_cf(omega, S0, r, q, T, kappa, theta, eta, rho, v0):
2     x = np.log(S0)
3     xi = kappa - 1j * rho * eta * omega
4     d = np.sqrt(xi**2 + eta**2 * (omega**2 + 1j * omega))
5     g = (xi - d) / (xi + d)
6
7     C = ((r - q) * 1j * omega * T
8          + (kappa * theta / eta**2) * (
9              (xi - d) * T - 2 * np.log((1 - g * np.exp(-d * T)) / (1 - g))
10          ))
11     D = ((xi - d) / eta**2) * (1 - np.exp(-d * T)) / (1 - g * np.exp(-d * T))
```

```

12
13     return np.exp(C + D * v0 + 1j * omega * x)

```

and finally, the FFT implementation of Carr-Madan using Simpson weights:

```

1  def carr_madan_fft(S0, r, q, T, kappa, theta, eta, rho, v0,
2                      alpha=1.5, N=4096, eta_grid=0.25):
3      lam = 2 * np.pi / (N * eta_grid)
4      b   = 0.5 * N * lam
5
6      j   = np.arange(1, N + 1)
7      v_j = eta_grid * (j - 1)
8
9      cf  = lambda omega: heston_cf(omega, S0, r, q, T, kappa, theta, eta, rho, v0)
10
11     psi = (np.exp(-r * T) * cf(v_j - (alpha + 1) * 1j)
12           / (alpha**2 + alpha - v_j**2 + 1j * (2 * alpha + 1) * v_j))
13
14     w   = (eta_grid / 3) * (3 + (-1)**j - (j == 1).astype(float))
15     k_u = -b + lam * (np.arange(1, N + 1) - 1)
16
17     x    = np.fft.fft(np.exp(1j * b * v_j) * psi * w)
18     calls = (np.exp(-alpha * k_u) / np.pi) * np.real(x)
19     strikes = np.exp(k_u)
20
21     return strikes, calls

```

Let's compare MC to Carr-Mada:

```

1  # — parameters —————
2
3  params = dict(S0=100, r=0.05, q=0, T=1.0,
4               kappa=2.0, theta=0.04, eta=0.3, rho=-0.7, v0=0.04)
5
6  strikes_to_price = np.arange(60, 141, 5, dtype=float)
7
8  # — MC ground truth —————
9
10 t0 = time.perf_counter()
11 mc_prices = np.array([heston_mc(**params, K=K) for K in strikes_to_price])
12 mc_time   = time.perf_counter() - t0
13
14 # — Carr-Madan —————

```



```

15
16 t0 = time.perf_counter()
17 cm_strikes, cm_calls = carr_madan_fft(**params)
18 cm_time = time.perf_counter() - t0
19 cm_prices = np.interp(strikes_to_price, cm_strikes, cm_calls)
20
21 # — print results —————
22
23 print(f"{'Strike':>8} {'MC':>10} {'Carr-Madan':>12} {'Error':>10}")
24 print("-" * 44)
25 for K, mc, cm in zip(strikes_to_price, mc_prices, cm_prices):
26     print(f"{K:>8.0f} {mc:>10.4f} {cm:>12.4f} {abs(mc-cm):>10.4f}")
27
28 print(f"\nMC time:           {mc_time:.2f}s ({len(strikes_to_price)} strikes, 100k paths)
29 print(f"Carr-Madan time: {cm_time*1000:.2f}ms (all strikes simultaneously)")
30
31 # — plot —————
32
33 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
34
35 ax1.plot(strikes_to_price, mc_prices, 'o', label='Monte Carlo', color='steelblue')
36 ax1.plot(strikes_to_price, cm_prices, '-', label='Carr-Madan', color='tomato', linewidth=
37 ax1.set_xlabel('Strike')
38 ax1.set_ylabel('Call Price')
39 ax1.set_title('Heston Call Prices')
40 ax1.legend()
41 ax1.grid(True, alpha=0.3)
42
43 ax2.bar(['Monte Carlo\n(100k paths, 17 strikes)', 'Carr-Madan FFT\n(all strikes)'],
44         [mc_time, cm_time],
45         color=['steelblue', 'tomato'])
46 ax2.set_ylabel('Time (seconds)')
47 ax2.set_title('Computation Time')
48 ax2.grid(True, alpha=0.3, axis='y')
49
50 plt.tight_layout()
51 plt.show()

```

1	Strike	MC	Carr-Madan	Error
2	-----			
3	60	43.0926	43.0487	0.0439
4	65	38.4447	38.4019	0.0428
5	70	33.8703	33.8281	0.0423
6	75	29.4033	29.3614	0.0420
7	80	25.0861	25.0447	0.0415

8	85	20.9701	20.9298	0.0403
9	90	17.1172	17.0758	0.0415
10	95	13.5838	13.5438	0.0401
11	100	10.4286	10.3950	0.0336
12	105	7.7081	7.6798	0.0283
13	110	5.4500	5.4306	0.0194
14	115	3.6676	3.6572	0.0104
15	120	2.3451	2.3334	0.0117
16	125	1.4146	1.4061	0.0085
17	130	0.8051	0.8002	0.0049
18	135	0.4343	0.4309	0.0034
19	140	0.2238	0.2206	0.0031
20				
21	MC time: 6.46s (17 strikes, 100k paths each)			
22	Carr-Madan time: 1.05ms (all strikes simultaneously)			



The results speak for themselves! Even at 100k paths, which isn't all that much, and gives somewhat inaccurate results, running MC takes around 6.5 seconds, while Carr-Madan only takes a millisecond. But we can do even better!

Lewis

Lewis (2001) reformulates option pricing as a complex contour integral. The result is a cleaner formula than Carr-Madan; no dampening parameter alpha to tune, and the si

condition is easily satisfied for essentially all models used in practice.

The core idea

The option price is a discounted expectation under Q (We are restating this since the paper has slightly different notation):

$$V(S_0) = e^{-rT} \int_{-\infty}^{\infty} w(x) q_T(x) dx$$

where $x = \log S_T$, $w(x)$ is the payoff function, and $q_T(x)$ is the risk-neutral density of $\log S_T$. This is an inner product of w and q_T . Parseval's theorem says inner products are preserved under Fourier transformation:

$$\int_{-\infty}^{\infty} w(x) \overline{q_T(x)} dx = \frac{1}{2\pi} \int_{-\infty}^{\infty} \hat{w}(\xi) \overline{\hat{q}_T(\xi)} d\xi$$

Since $q_T(x)$ is real, we have

$$\overline{\hat{q}_T(\xi)} = \hat{q}_T(-\xi) = \phi_T(-\xi)$$

So:

$$V(S_0) = \frac{e^{-rT}}{2\pi} \int_{-\infty}^{\infty} \hat{w}(\xi) \phi_T(-\xi) d\xi$$

The problem

That formula looks clean, but there's a catch. For a call, the payoff is

$$w(x) = (e^x - K)^+,$$

which grows like e^x as $x \rightarrow \infty$. Its ordinary Fourier transform diverges!

Carr-Madan's fix was to multiply the payoff by a dampening factor before transforming. Lewis takes a cleaner route: allow the transform variable to be complex.

Generalized Fourier Transform

Let $z = u+iv$ be a complex number. Define the generalized Fourier transform of the payoff:

$$\hat{w}(z) = \int_{-\infty}^{\infty} e^{izx} w(x) dx$$

Now $e^{izx} = e^{iux - vx}$, and the factor e^{-vx} provides the decay that kills the growth of $w(x)$, as long as v is large enough. For the call:

$$\hat{w}(z) = \int_{\ln K}^{\infty} e^{izx} (e^x - K) dx$$

This does not converge unless $\text{Im}(z) > 1$. With $\text{Im}(z) > 1$ we have:

$$\hat{w}(z) = -\frac{K^{iz+1}}{z^2 - iz}, \quad \Im(z) > 1$$

We call this area the **strip of regularity** for the call payoff:

$$\mathcal{S}_w = \{z : \Im(z) > 1\}$$

You can refer to the following table from the paper for other options:

Table 3.1
Generalized Fourier Transforms for Various Financial Claims

Financial Claim (Option)	Payoff Function: $w(x)$	Payoff Transform: $\hat{w}(z) = \mathcal{F}[w(x)]$	Strip of regularity $\mathcal{S}_w$
Call option	$(e^x - K)^+$	$-\frac{K^{iz+1}}{z^2 - iz}$	$\text{Im } z > 1$
Put option	$(K - e^x)^+$	$-\frac{K^{iz+1}}{z^2 - iz}$	$\text{Im } z < 0$
Covered call or cash-secured put	$\min(e^x, K)$	$\frac{K^{iz+1}}{z^2 - iz}$	$0 < \text{Im } z < 1$
Delta function	$\delta(x - \ln K)$	K^{iz}	Entire z -plane
Money market	1	$2\pi\delta(z)$	$\text{Im } z = 0$

Notes: The table uses $x = \log S_T$, where S_T is the security price at expiration, and $x^+ = \max(x, 0)$. The parameter K is a strike (exercise) price. The expression $\delta(x - x_0)$ is Dirac delta function. The symbol $\mathcal{F}[\cdot]$ indicates taking a Fourier transform.

The characteristic function has its own regularity strip. For the characteristic function of X_T to exist, we need

$$\mathbb{E}[e^{-vX_T}] < \infty.$$

The set of valid complex v with this property forms the strip $\mathcal{S}_X$.

Define the **reflected strip**

$$\mathcal{S}_X^* = \{z : -z \in \mathcal{S}_X\}.$$

This is where $\phi_T(-z)$ is regular, which is what appears in the pricing formula.

Shifting the contour

By Cauchy's theorem, the integration contour may be shifted to any horizontal line inside the **common strip of regularity**:

$$\int_{-\infty}^{\infty} \hat{w}(\xi) \phi_T(-\xi) d\xi = \int_{iv-\infty}^{iv+\infty} \hat{w}(\xi) \phi_T(-\xi) d\xi, \quad v \in \mathcal{S}_V = \mathcal{S}_w \cap \mathcal{S}_X^*$$

The Lewis Formula

The option price finally becomes:

$$V(S_0) = \frac{e^{-rT}}{2\pi} \int_{iv-\infty}^{iv+\infty} e^{-izY} \phi_T(-z) \hat{w}(z) dz$$

where $v = \text{Im}(z)$ in $\mathcal{S}_V$ and $Y = \ln S_0 + (r-q)T$ is the log forward price, and ϕ denotes the characteristic function of the centred log-return $X_T = \log S_T - Y$.

Call Pricing Formula

Substituting the call payoff transform $\hat{w}$ from above:

$$C(S_0) = \frac{e^{-rT}}{2\pi} \int_{iv-\infty}^{iv+\infty} e^{-izY} \phi_T(-z) \cdot \left(-\frac{K^{iz+1}}{z^2 - iz}\right) dz, \quad v \in (1, b)$$

where b is some model-dependent upper bound below which the characteristic function stays regular. Lewis's preferred contour is $v=1/2$, obtained by shifting the contour from $(1, b)$ down to $(0, 1)$ via the residue theorem and picking up the pole at $z=i$. That gives

$$C(S_0) = S_0 e^{-qT} - \frac{\sqrt{S_0 K} e^{-(r+q)T/2}}{\pi} \int_0^\infty \Re[e^{iuk} \phi_T(u - \frac{i}{2})] \frac{du}{u^2 + \frac{1}{4}},$$

where $k = \ln(S_0/K) + (r-q)T$ is the log-moneyness. The contour $v=1/2$ sits symmetrically between the poles at $z=0$ and $z=i$, and requires only

$$\mathbb{E}^\mathbb{Q}[S_T^{1/2}] < \infty,$$

satisfied by practically every standard model, and is a much weaker requirement than

$$\mathbb{E}^{\mathbb{Q}}[S_T^{\alpha+1}] < \infty,$$

that Carr-Madan needs for $\alpha > 0$.

Python Implementation

Lewis uses the characteristic function of X_T :

```

1  def heston_cf_xt(omega, S0, r, q, T, kappa, theta, eta, rho, v0):
2      """CF of  $X_T = \log(S_T/S_0) - (r-q)T$  under Heston."""
3      xi = kappa - 1j * rho * eta * omega
4      d = np.sqrt(xi**2 + eta**2 * (omega**2 + 1j * omega))
5      g = (xi - d) / (xi + d)
6
7      C = (kappa * theta / eta**2) * (
8          (xi - d) * T - 2 * np.log((1 - g * np.exp(-d * T)) / (1 - g))
9      )
10     D = ((xi - d) / eta**2) * (1 - np.exp(-d * T)) / (1 - g * np.exp(-d * T))
11
12     # no (r-q) drift term since  $X_T$  is already centered
13     return np.exp(C + D * v0 - 0.5 * 1j * omega * eta**2 * 0)

```

Now since Lewis doesn't use FFT by design, we have to compute strikes one by one using the contour integral, or we are smart and vectorize:

```

1  def lewis_calls_vectorized(S0, strikes, r, q, T, kappa, theta, eta, rho, v0,
2                             N=4096, eta_grid=0.25):
3      u_j = eta_grid * np.arange(0, N)
4
5      # Simpson weights
6      w = np.ones(N) * 2
7      w[1::2] = 4
8      w[0] = w[-1] = 1
9      w *= eta_grid / 3
10
11     phi = heston_cf_xt(u_j - 0.5j, S0, r, q, T, kappa, theta, eta, rho, v0)
12     phi /= (u_j**2 + 0.25)
13
14     # k for each strike: shape (n_strikes,)
15     k = np.log(S0 / strikes) + (r - q) * T
16

```

```

17     # outer product: shape (n_strikes, N)
18     phase = np.exp(1j * np.outer(k, u_j))
19
20     I = np.real(phase * phi[np.newaxis, :]) @ w / np.pi
21
22     return S0 * np.exp(-q * T) - np.sqrt(S0 * strikes) * np.exp(-(r + q) * T / 2) * I

```



And the comparison:

```

1  t0 = time.perf_counter()
2  lw_prices = lewis_calls_vectorized(strikes=strikes_to_price, **params)
3  lw_time = time.perf_counter() - t0
4
5  print(f"{'Strike':>8} {'MC':>10} {'Carr-Madan':>12} {'Lewis':>10} "
6        f"{'CM Err':>10} {'LW Err':>10}")
7  print("-" * 64)
8  for K, mc, cm, lw in zip(strikes_to_price, mc_prices, cm_prices, lw_prices):
9      print(f"{K:>8.0f} {mc:>10.4f} {cm:>12.4f} {lw:>10.4f} "
10            f"{abs(mc-cm):>10.4f} {abs(mc-lw):>10.4f}")
11
12  print(f"\nMC time:           {mc_time:.2f}s")
13  print(f"Carr-Madan time:    {cm_time*1000:.2f}ms")
14  print(f"Lewis time:         {lw_time*1000:.2f}ms")
15
16  fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
17
18  ax1.plot(strikes_to_price, mc_prices, 'o', label='Monte Carlo', color='steelblue')
19  ax1.plot(strikes_to_price, cm_prices, '-', label='Carr-Madan', color='tomato', linewidth=
20  ax1.plot(strikes_to_price, lw_prices, '--', label='Lewis', color='seagreen', linewidth=2)
21  ax1.set_xlabel('Strike')
22  ax1.set_ylabel('Call Price')
23  ax1.set_title('Heston Call Prices')
24  ax1.legend()
25  ax1.grid(True, alpha=0.3)
26
27  ax2.bar(['Carr-Madan\nFFT', 'Lewis\nVectorized'],
28          [cm_time, lw_time],
29          color=['tomato', 'seagreen'])
30  ax2.set_ylabel('Time (seconds)')
31  ax2.set_title('Computation Time (excl. Monte Carlo)')
32  ax2.grid(True, alpha=0.3, axis='y')
33
34  plt.tight_layout()
35  plt.show()

```

1	Strike	MC	Carr-Madan	Lewis	CM Err	LW Err
2	-----					
3	60	43.0926	43.0487	43.1459	0.0439	0.0533
4	65	38.4447	38.4019	38.5020	0.0428	0.0573
5	70	33.8703	33.8281	33.9311	0.0423	0.0608
6	75	29.4033	29.3614	29.4673	0.0420	0.0640
7	80	25.0861	25.0447	25.1536	0.0415	0.0674
8	85	20.9701	20.9298	21.0417	0.0403	0.0716
9	90	17.1172	17.0758	17.1902	0.0415	0.0730
10	95	13.5838	13.5438	13.6612	0.0401	0.0774
11	100	10.4286	10.3950	10.5150	0.0336	0.0864
12	105	7.7081	7.6798	7.8025	0.0283	0.0944
13	110	5.4500	5.4306	5.5570	0.0194	0.1070
14	115	3.6676	3.6572	3.7858	0.0104	0.1182
15	120	2.3451	2.3334	2.4652	0.0117	0.1201
16	125	1.4146	1.4061	1.5412	0.0085	0.1266
17	130	0.8051	0.8002	0.9381	0.0049	0.1329
18	135	0.4343	0.4309	0.5717	0.0034	0.1375
19	140	0.2238	0.2206	0.3646	0.0031	0.1409
20						
21	MC time:	8.09s				
22	Carr-Madan time:	1.18ms				
23	Lewis time:	2.49ms				



As you can see, when using the same N and η as with Carr-Madan, Lewis starts diverging pretty quickly and is still over twice as slow. If we increase N and reduce η , we are quickly multiple times slower than Carr-Madan when pricing many strikes.

once. On the other hand, we get rid of the dampening parameter alpha, which can cause real problems.

Let's finally look at our method of choice, which performs better than both Lewis and Carr-Madan.

The COS Method

Both Carr-Madan and Lewis price options by evaluating a Fourier integral numerically. The integrand is oscillatory, which is expensive to integrate accurately with equally spaced quadrature.

The COS method (2008) takes a different route. Instead of integrating the characteristic function directly, it expands the risk-neutral density in a cosine series and reads the series coefficients directly off the characteristic function, with those coefficients often being available analytically.

Fourier-Cosine Series Expansion

Starting from risk-neutral pricing (once again, slightly different notation!):

$$v(x, 0) = e^{-rT} \int_{-\infty}^{\infty} v(y, T) f(y|x) dy$$

where $x = \ln(S_0/K)$, $y = \ln(S_T/K)$, $v(y, T)$ is the payoff, and $f(y|x)$ is the risk-neutral transition density of the log-return.

We truncate to a finite interval $[a, b]$:

$$v(x, 0) \approx e^{-rT} \int_a^b v(y, T) f(y|x) dy$$

Now expand $f(y|x)$ in a cosine series on $[a, b]$:

$$f(y|x) \approx \sum_{k=0}^{N-1} 'A_k(x) \cos(k\pi \frac{y-a}{b-a})$$

where the ' in the sum means the $k=0$ term is weighted by 1/2, and the cosine coefficients are:

$$A_k(x) = \frac{2}{b-a} \int_a^b f(y|x) \cos(k\pi \frac{y-a}{b-a}) dy$$

Substituting into the pricing integral and swapping sum and integral (which we can thanks to Fubini):

$$v(x, 0) \approx \frac{b-a}{2} e^{-rT} \sum_{k=0}^{N-1} 'A_k(x) V_k$$

where:

$$V_k = \frac{2}{b-a} \int_a^b v(y, T) \cos(k\pi \frac{y-a}{b-a}) dy$$

Note that the V_k are the cosine series coefficients of $v(y, T)$ in y . Thus we have transformed the product of two real functions, $f(y|x)$ and $v(y, T)$, to that of their Fourier cosine series coefficients.

Coefficients from the characteristic function

The density $f(y|x)$ is unknown, that's the whole reason we use Fourier methods in the first place. Comparing the definition of $A_k(x)$ with the characteristic function $\phi_T(w)$, one finds that:

$$A_k(x) = \frac{2}{b-a} \Re[\phi_1(\frac{k\pi}{b-a}) \exp(-i \frac{k\pi a}{b-a})]$$

where ϕ_1 is the characteristic function computed over the truncated range $[a, b]$. Since $[a, b]$ was chosen to capture essentially all the density mass, ϕ_1 is approximately ϕ_T and so:

$$A_k(x) \approx F_k = \frac{2}{b-a} \Re[\phi_T(\frac{k\pi}{b-a}) \exp(-i \frac{k\pi a}{b-a})]$$

The COS Formula

Finally, replacing A_k by F_k in the formula for $V(x, t_0)$, we obtain:

$$v(x, 0) \approx e^{-rT} \sum_{k=0}^{N-1} \Re[\phi_T(\frac{k\pi}{b-a}) e^{-ik\pi a/(b-a)}] V_k$$

This is the COS formula.

Payoff coefficients for calls and puts

For a call, $v(y, T) = K(e^y - 1)^+$, the integration domain reduces to $[0, b]$ since the payoff is zero for $y < 0$. The coefficients are:

$$V_k^{\text{call}} = \frac{2}{b-a} K(\chi_k(0, b) - \psi_k(0, b))$$

where χ and ψ have closed forms:

$$\begin{aligned} \chi_k(c, d) := & \frac{1}{1 + \left(\frac{k\pi}{b-a}\right)^2} \left[\cos\left(k\pi \frac{d-a}{b-a}\right) e^d - \cos\left(k\pi \frac{c-a}{b-a}\right) e^c \right. \\ & \left. + \frac{k\pi}{b-a} \sin\left(k\pi \frac{d-a}{b-a}\right) e^d - \frac{k\pi}{b-a} \sin\left(k\pi \frac{c-a}{b-a}\right) e^c \right] \end{aligned}$$

and

$$\psi_k(c, d) := \begin{cases} \left[\sin\left(k\pi \frac{d-a}{b-a}\right) - \sin\left(k\pi \frac{c-a}{b-a}\right) \right] \frac{b-a}{k\pi} & k \neq 0, \\ (d - c) & k = 0. \end{cases}$$

I couldn't be bothered typing this out in LaTeX

For a put the integration domain is $[a, 0]$, and

$$V_k^{\text{put}} = \frac{2}{b-a} K(-\chi_k(a, 0) + \psi_k(a, 0))$$

In practice, you can price puts or calls, and recover the other via put-call parity

$$C = P + S_0 e^{-qT} - K e^{-rT}.$$

Truncation Range

The interval $[a, b]$ must contain enough of the density without being so wide that the cosine approximation needs excessive terms to resolve the tails. Fang & Oosterlee propose using cumulants of $\ln(S_T/K)$:

$$[a, b] = [c_1 - L\sqrt{c_2 + \sqrt{c_4}}, c_1 + L\sqrt{c_2 + \sqrt{c_4}}], \quad L = 10$$

where c_1, c_2, c_4 are the first, second and fourth cumulants of $\ln(S_T/K)$. Including c_4 matters for short maturities and fat-tailed processes where the density is sharply peaked.

You can additionally also include a c_6 term as well for extremely short maturities, the sixth cumulant can be difficult to derive for many models. You can also go the other direction and only include the c_1 and c_2 terms.

Benefits of the COS method

For Carr-Madan and Lewis, accuracy is limited by how well a trapezoidal or Simpson's sum approximates a slowly decaying oscillatory integral. The COS method avoids this entirely. For smooth densities like GBM, Heston, and many Lévy processes at moderate maturities, the density is C^∞ on $[a, b]$ and the cosine coefficients decay exponentially, yielding exponential convergence.

The computational complexity is $O(N)$ for a single strike. For a vector of strikes, the F_k are the same across all strikes, and the V_k differ only through the x -dependent phase; pricing M strikes simultaneously costs $O(MN)$ via matrix-vector multiplication.

Python Implementation

We need the first and second cumulants of $\log(S_T/K)$ under Heston:

```

1  def heston_cumulants(r, q, T, kappa, theta, eta, rho, v0):
2      """First and second cumulants of log(S_T/K) under Heston."""
3      c1 = (r - q) * T + (1 - np.exp(-kappa * T)) * (theta - v0) / (2 * kappa) - 0.5 * thet
4
5      c2 = (1 / (8 * kappa**3)) * (
6          eta * T * kappa * np.exp(-kappa * T) * (v0 - theta) * (8 * kappa * rho - 4 * eta)
7          + kappa * rho * eta * (1 - np.exp(-kappa * T)) * (16 * theta - 8 * v0)
8          + 2 * theta * kappa * T * (-4 * kappa * rho * eta + eta**2 + 4 * kappa**2)
9          + eta**2 * ((theta - 2 * v0) * np.exp(-2 * kappa * T)
10             + theta * (6 * np.exp(-kappa * T) - 7) + 2 * v0)
11          + 8 * kappa**2 * (v0 - theta) * (1 - np.exp(-kappa * T))
12      )
13
14      return c1, c2

```

Fast vectorized implementation of the COS method:

```

1  def cos_calls(S0, strikes, r, q, T, kappa, theta, eta, rho, v0, N=128, L=12):
2      c1, c2 = heston_cumulants(r, q, T, kappa, theta, eta, rho, v0)
3      a = c1 - L * np.sqrt(abs(c2))
4      b = c1 + L * np.sqrt(abs(c2))
5
6      k = np.arange(N)
7      omega = k * np.pi / (b - a)
8
9      def varphi(omega):
10         xi = kappa - 1j * rho * eta * omega
11         d = np.sqrt(xi**2 + eta**2 * (omega**2 + 1j * omega))
12         g = (xi - d) / (xi + d)
13         C = ((r - q) * 1j * omega * T
14             + (kappa * theta / eta**2) * (
15                 (xi - d) * T - 2 * np.log((1 - g * np.exp(-d * T)) / (1 - g))
16             ))
17         D = ((xi - d) / eta**2) * (1 - np.exp(-d * T)) / (1 - g * np.exp(-d * T))
18         return np.exp(C + D * v0)
19
20     def chi(c, d):
21         kpi = np.where(k != 0, omega, 1.0)
22         val = (np.cos(kpi * (d - a)) * np.exp(d)
23             - np.cos(kpi * (c - a)) * np.exp(c)
24             + kpi * np.sin(kpi * (d - a)) * np.exp(d)
25             - kpi * np.sin(kpi * (c - a)) * np.exp(c)) / (1 + kpi**2)
26         val[0] = np.exp(d) - np.exp(c)

```

```

27         return val
28
29     def psi(c, d):
30         val = np.empty(N)
31         val[0] = d - c
32         val[1:] = (b - a) / (k[1:] * np.pi) * (
33             np.sin(omega[1:] * (d - a)) - np.sin(omega[1:] * (c - a))
34         )
35         return val
36
37     # precompute strike-independent quantities
38     basis = varphi(omega) * np.exp(-1j * omega * a) # vphi * e^{-i*omega*a}
39     A = np.real(basis)
40     B = np.imag(basis)
41     Uk = (2 / (b - a)) * (-chi(a, 0) + psi(a, 0))
42
43     # cache constants
44     disc = np.exp(-r * T)
45     fwd_S0 = S0 * np.exp(-q * T)
46     disc_K = np.exp(-r * T)
47
48     # vectorize over strikes
49     x = np.log(S0 / strikes) # shape (M,)
50     theta_mx = np.outer(x, omega) # shape (M, N)
51     F = A * np.cos(theta_mx) - B * np.sin(theta_mx) # shape (M, N)
52     F[:, 0] *= 0.5
53
54     puts = disc * strikes * (F @ Uk)
55     calls = puts + fwd_S0 - strikes * disc_K
56     return calls

```

And last but not least, the comparison code:

```

1  t0 = time.perf_counter()
2  cos_prices = cos_calls(strikes=strikes_to_price, **params)
3  cos_time = time.perf_counter() - t0
4
5  print(f"{'Strike':>8} {'MC':>10} {'Carr-Madan':>12} {'Lewis':>10} {'COS':>10} "
6        f"{'CM Err':>10} {'LW Err':>10} {'COS Err':>10}")
7  print("-" * 84)
8  for K, mc, cm, lw, cos in zip(strikes_to_price, mc_prices, cm_prices, lw_prices, cos_prices):
9      print(f"{K:>8.0f} {mc:>10.4f} {cm:>12.4f} {lw:>10.4f} {cos:>10.4f} "
10            f"{abs(mc-cm):>10.4f} {abs(mc-lw):>10.4f} {abs(mc-cos):>10.4f}")
11
12  print(f"\nMC time: {mc_time:.2f}s")

```


```

13 print(f"Carr-Madan time:  {cm_time*1000:.2f}ms")
14 print(f"Lewis time:      {lw_time*1000:.2f}ms")
15 print(f"COS time:        {cos_time*1000:.2f}ms")
16
17 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
18
19 ax1.plot(strikes_to_price, mc_prices, 'o', label='Monte Carlo', color='steelblue')
20 ax1.plot(strikes_to_price, cm_prices, '-', label='Carr-Madan', color='tomato', linewidth=
21 ax1.plot(strikes_to_price, lw_prices, '--', label='Lewis', color='seagreen', linewidth=2)
22 ax1.plot(strikes_to_price, cos_prices, ':', label='COS', color='purple', linewidth=2)
23 ax1.set_xlabel('Strike')
24 ax1.set_ylabel('Call Price')
25 ax1.set_title('Heston Call Prices')
26 ax1.legend()
27 ax1.grid(True, alpha=0.3)
28
29 ax2.bar(['Carr-Madan\nFFT', 'Lewis\nVectorized', 'COS'],
30         [cm_time, lw_time, cos_time],
31         color=['tomato', 'seagreen', 'purple'])
32 ax2.set_ylabel('Time (seconds)')
33 ax2.set_title('Computation Time (excl. Monte Carlo)')
34 ax2.grid(True, alpha=0.3, axis='y')
35
36 plt.tight_layout()
37 plt.show()

```

	Strike	MC	Carr-Madan	Lewis	COS	CM Err	LW Err	COS Err
1								
2								
3	60	43.0926	43.0487	43.1459	43.0487	0.0439	0.0533	0.0439
4	65	38.4447	38.4019	38.5020	38.4019	0.0428	0.0573	0.0428
5	70	33.8703	33.8281	33.9311	33.8280	0.0423	0.0608	0.0424
6	75	29.4033	29.3614	29.4673	29.3612	0.0420	0.0640	0.0421
7	80	25.0861	25.0447	25.1536	25.0446	0.0415	0.0674	0.0416
8	85	20.9701	20.9298	21.0417	20.9297	0.0403	0.0716	0.0403
9	90	17.1172	17.0758	17.1902	17.0753	0.0415	0.0730	0.0419
10	95	13.5838	13.5438	13.6612	13.5434	0.0401	0.0774	0.0404
11	100	10.4286	10.3950	10.5150	10.3942	0.0336	0.0864	0.0344
12	105	7.7081	7.6798	7.8025	7.6788	0.0283	0.0944	0.0293
13	110	5.4500	5.4306	5.5570	5.4303	0.0194	0.1070	0.0197
14	115	3.6676	3.6572	3.7858	3.6562	0.0104	0.1182	0.0114
15	120	2.3451	2.3334	2.4652	2.3326	0.0117	0.1201	0.0125
16	125	1.4146	1.4061	1.5412	1.4057	0.0085	0.1266	0.0089
17	130	0.8051	0.8002	0.9381	0.7996	0.0049	0.1329	0.0055
18	135	0.4343	0.4309	0.5717	0.4303	0.0034	0.1375	0.0039

19	140	0.2238	0.2206	0.3646	0.2203	0.0031	0.1409	0.0035
20								
21	MC time:	8.09s						
22	Carr-Madan time:	1.18ms						
23	Lewis time:	2.49ms						
24	COS time:	0.34ms						



Awesome! We are more than 3 times faster than Carr-Madan, while getting rid of annoying parameters and gaining exponential convergence!


The method naturally vectorizes into dense linear algebra operations, meaning you could still make this way way faster if you optimized it properly.

Conclusion

Apologies for the delay since the last article! This one took a little longer than I expected since I had to verify all the math, and turn 3 papers into one article instead of the usual one paper.

I hope you enjoyed this rather math-heavy article anyway! I plan on covering more options content soon.


Other Articles You Would Enjoy



How to Trade the Volatility Smile

VERTOX • JUNE 29, 2025


[Read full story](#) →



An Introduction to Option Market Making

VERTOX • SEPTEMBER 19, 2025

[Read full story](#) →



Pricing Illiquid Options Using Proxy Volatility Surfaces

VERTOX • NOVEMBER 24, 2025

[Read full story](#) →

Quant Corner

Join Quant Corner, a community that actually *does things*. Tournaments, live Q&As, open discussions, and the best place to shape what gets built next.

Join here: <https://discord.gg/X7TsxKNbXg>

Subscribe to VertoxQuant

Hundreds of paid subscribers

Applied quantitative research on trading, risk, and systematic strategy design.

[Upgrade to paid](#)



19 Likes · 5 Restacks

Discussion about this post

Comments

Restacks



Write a comment...



Sacha Savic  2d

 Liked by Vertox

So interesting. I was wondering about using FFT to price options after a Signals physics last week rock!

 LIKE (1)  REPLY



1 reply by Vertox

1 more comment...